

a-

The nested loops inside the `add()` function should be parallelized because they are performing the same operation on different elements of the matrices.

b-

The loops in the `add()` function have no loop-carried dependencies since each element of the output matrix "C" is computed independently using only the corresponding elements of the input matrices "A" and "B".

c-

Here's the parallel version of the program using OpenMP:

....

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
#define N 16
```

```
void add(int A[][N], int B[][N], int C[][N])
```

```
{
```

```
    int i, j;
```

```
    #pragma omp parallel for private(i, j)
```

```
    for (i = 0; i < N; i++)
```

```
    {
```

```
        for (j = 0; j < N; j++)
```

```
        {
```

```
            C[i][j] = A[i][j] + B[i][j];
```

```
        }
```

```
    }
```

```

    }
    int main ( )
    {
        int C[N][N];
        int A[N][N] = {(1,1,1,1....), (2,2,2,2....),
                        (3,3,3,3....), --- (4,4,4,4....)};

        add(A,B,C);
        printf("Result Matrix is \n");
        int i,j;
        for(i=0; i<N; i++)
        {
            for(j=0; j<N; j++)
            {
                printf("%d", C[i][j]);
            }
            printf("\n");
        }
        return 0;
    }

```